

```
In [1]: import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

```
In [2]: df=pd.read_csv("C:\\Users\\JOHNSON\\Downloads\\my notes\\linear\\assignment\\As
df
```

```
Out[2]:
```

	Gender	Height	Weight
0	Male	73.847017	241.893563
1	Male	68.781904	162.310473
2	Male	74.110105	212.740856
3	Male	71.730978	220.042470
4	Male	69.881796	206.349801
...	...	...	...
9995	Female	66.172652	136.777454
9996	Female	67.067155	170.867906
9997	Female	63.867992	128.475319
9998	Female	69.034243	163.852461
9999	Female	61.944246	113.649103

10000 rows × 3 columns

```
In [3]: df.head
```

```
Out[3]: <bound method NDFrame.head of
0      Male  73.847017  241.893563
1      Male  68.781904  162.310473
2      Male  74.110105  212.740856
3      Male  71.730978  220.042470
4      Male  69.881796  206.349801
...      ...      ...      ...
9995  Female  66.172652  136.777454
9996  Female  67.067155  170.867906
9997  Female  63.867992  128.475319
9998  Female  69.034243  163.852461
9999  Female  61.944246  113.649103
```

[10000 rows x 3 columns]>

In [4]: df.tail

```
Out[4]: <bound method NDFrame.tail of      Gender      Height      Weight
0      Male  73.847017  241.893563
1      Male  68.781904  162.310473
2      Male  74.110105  212.740856
3      Male  71.730978  220.042470
4      Male  69.881796  206.349801
...      ...      ...      ...
9995  Female  66.172652  136.777454
9996  Female  67.067155  170.867906
9997  Female  63.867992  128.475319
9998  Female  69.034243  163.852461
9999  Female  61.944246  113.649103

[10000 rows x 3 columns]>
```

In [6]: x=df['Height'].array.reshape(-1,1)  
x

```
Out[6]: <PandasArray>
[
  [73.847017017515],
  [68.7819040458903],
  [74.1101053917849],
  [71.7309784033377],
  [69.8817958611153],
  [67.2530156878065],
  [68.7850812516616],
  [68.3485155115879],
  [67.018949662883],
  [63.4564939783664],
  [71.1953822829745],
  [71.6408051192206],
  [64.7663291334055],
  [69.2830700967204],
  [69.2437322298112],
  [67.6456197004212],
  [72.4183166259878],
  [62.0743357310611],
  ...
]
```

In [8]: y=df['Weight']  
y

```
Out[8]: 0      241.893563
1      162.310473
2      212.740856
3      220.042470
4      206.349801
...
9995    136.777454
9996    170.867906
9997    128.475319
9998    163.852461
9999    113.649103
Name: Weight, Length: 10000, dtype: float64
```

```
In [9]: #splitting the data
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
In [10]: model=LinearRegression()
model.fit(x_train,y_train)
```

Out[10]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [11]: score=model.score(x_train,y_train)
score
```

Out[11]: 0.8545053200432668

```
In [12]: model.coef_
```

Out[12]: array([7.70218561])

```
In [13]: model.intercept_
```

Out[13]: -349.78782058244576

## Optimisation

```
In [14]: #Defining parameters
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error
```

```
In [17]: param_grid={
    'copy_X':[True,False],
    'fit_intercept':[True,False],
    'n_jobs':[True,False],
    'positive':[True,False]
}
```

```
In [18]: #performing grid with cross_validation
grid_search=GridSearchCV(model,param_grid,cv=5,scoring='neg_mean_squared_error')
grid_search.fit(x_train,y_train)
```

```
Out[18]: GridSearchCV(cv=5, estimator=LinearRegression(),
                    param_grid={'copy_X': [True, False],
                                'fit_intercept': [True, False],
                                'n_jobs': [True, False], 'positive': [True, False]},
                    scoring='neg_mean_squared_error')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [19]: #Best model
best_model=grid_search.best_estimator_
best_model
```

```
Out[19]: LinearRegression(n_jobs=True, positive=True)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [20]: best_score=model.score(x_test,y_test)
best_score
```

```
Out[20]: 0.85773177770385
```

```
In [21]: y_pred=best_model.predict(x_test)
y_pred
```

```
Out[21]: array([179.25399046, 180.34848321, 161.62288801, ..., 129.20288223,
                166.78470522, 101.81227499])
```

```
In [23]: MSE=mean_squared_error(y_test,y_pred)
MSE
```

```
Out[23]: 149.00350418448116
```

```
In [ ]:
```